

CSCI 4270 and 6270
Computational Vision,
Spring Semester 2025
Homework 3

Due: Wednesday, February 19, at 11:59:59 pm EST

This homework, worth **100 points** toward your overall homework grade, is a multistage problem of recognizing whether or not images are taken of the same scene and building a montage of images taken not only of the same scene, but from the same viewpoint (or taken of a planar surface). This problem will require a careful write-up describing your design decisions and results. Because of the extensive output needed, **do not** submit your solution as a Jupyter notebook, but instead submit a Python py file and separate files saving the text and some of the images output. See more details toward the end of this assignment.

Please note that there is significantly more to this assignment than just the code.

Overview

Here is the basic problem statement, which is elaborated on below: Given a folder of N images as input, for each pair of images I_i and I_j you must

1. Decide if I_i and I_j show the same scene.
2. If the decision for 1 is “yes”, then decide if I_i and I_j can be aligned accurately enough to form a mosaic.
3. If the decisions for both 1 and 2 are “yes”, then create and output a mosaic that aligns image I_i and I_j accurately.

It is possible that more than two images in a set of input images do show the same scene **and** may be combined into a mosaic. This is where the undergraduate and graduate versions of this assignment differ. In order to earn full credit, graduate students **must** produce a multi-image mosaic (in addition to the image pair mosaics). Undergraduates can earn a small amount of extra credit for multi-image mosaics. For graduate students this is the last 10 points on the assignment. For undergraduates, this is 5 points of extra credit. More on this below.

Details

Each image set you are given includes $N \geq 2$ images, I_1, \dots, I_N . Each image should be read in and processed as grayscale! For each pair of images, I_i and I_j , with $1 \leq i < j \leq N$, your code must do the following:

1. Extract the keypoints and descriptors in each image, using SIFT keypoints and descriptors. **Output** the number of keypoints in each image.
2. Match the keypoints and descriptors between the images. You may use `cv2.BFMatcher` or `cv2.FlannBasedMatcher` to do the matching. The decision about whether or not two keypoints match should be made using the descriptor ratio test as discussed in class. When a proposed match fails the ratio test it should be discarded; only keep those that pass the ratio test. Note that despite the ratio test, at this point there will often be errors in your keypoint matches.

Output:

- (a) The number of matches and the fraction of keypoints in each image that have a match. This fraction should be significantly less than 1.
 - (b) A single image showing I_i and I_j side-by-side with lines drawn between matched keypoints (see `cv2.drawMatches`). The lines should be different colors.
3. If the previous step produced too few matches overall or too small a percentage of matches, then stop attempting to match I_i and I_j — they show different scenes. (You will need to decide the criteria and the threshold or thresholds.) Otherwise proceed to the next step of matching. **Output** a message giving the decision made at this step.
4. Using the matches produced by keypoint description matching, use RANSAC to estimate the fundamental matrix $\mathbf{F}_{j,i}$ that maps pixel locations from I_i onto lines in I_j . Please review the significance of the fundamental matrix in your class notes! You may use `cv2.findFundamentalMat`. Do this with the `method` setting `cv2.FM_RANSAC`; you will have to explore the other parameter settings.

After estimating $\mathbf{F}_{j,i}$, you must determine which matches are “inliers” — consistent with the fundamental matrix. Specifically, if $\tilde{\mathbf{u}}_i$ (from image I_i) and $\tilde{\mathbf{u}}_j$ (from image I_j) are the homogeneous coordinate locations of a matching keypoint, then $\mathbf{a}_{j,i} = \mathbf{F}_{j,i}\tilde{\mathbf{u}}_i$ will be the coordinates describing the line in image j along which $\tilde{\mathbf{u}}_j$ should lie if it is a correct match. (This is the “epipolar line”.) While in theory $\tilde{\mathbf{u}}_j$ would be exactly on the line, in practice it may be slightly off. On the other hand, most incorrect matches will typically have $\tilde{\mathbf{u}}_j$ far from this line. Therefore you can determine which matches are inliers by measuring the distance between $\tilde{\mathbf{u}}_j$ and the line, and then you can count the number of keypoint matches that are within a small distance of the line. This is easy to do yourself as long as you determine the threshold and are careful to normalize $\mathbf{a}_{j,i}$ properly so that you can measure distances correctly. **However**, the `mask` array that `cv2.findFundamentalMat` returns does this for you! You are welcome to use it.

Output the following from this step:

- (a) The number and percentage of matches that are inliers.
 - (b) An image showing I_i and I_j side-by-side with lines drawn between the keypoints that form inlier matches (see `cv2.drawMatches`). Make each line a different color.
 - (c) An image showing the epipolar lines for the inlier matches drawn on image I_2 . Make each line a different color. This one may take a bit of work so we suggest saving it until after everything else is working.
5. If the previous step produced too few matches overall or too small a percentage of matches, then stop attempting to match I_i and I_j , and move on to the next image pair. (You will need to decide the criteria and the threshold or thresholds.) Otherwise proceed to the next step of matching. At this point your code will have made the decision that tells us whether or not I_i and I_j show the same scene. **Output** a message giving the decision made at this step.
6. Using the inlier matches from the fundamental matrix estimation step, estimate the parameters of the homography matrix $\mathbf{H}_{j,i}$ mapping I_i onto I_j . You may use `cv2.findHomography` and RANSAC. Using a criteria for deciding which matches are “inliers”, count the number of inliers for the homography matching between images.

Output the following from this step:

- (a) The number and percentage of inlier matches.
 - (b) An image showing I_i and I_j side-by-side with lines drawn between the keypoints that form inlier matches (see `cv2.drawMatches`). Make each line a different color.
7. Based on the number of inlier matches from fundamental matrix estimation and from homography estimation, decide whether or not the images can be accurately aligned. The decision should be “yes” if most of the inlier matches from the fundamental matrix estimate are also kept as inliers to the homography estimate. **Output** your decision and the reason for your decision.
 8. If the decision after the previous step is “yes” then build and output the mosaic of the two images. Try to come up with a relatively simple blending method that yields nice results instead of looking like one image is mapped and pasted on top of the other.

Multi-Image Mosaics

Here is a bit about forming multi-image mosaics, a problem you should leave until everything else is done. First, you need to remember which pairs of images can be aligned using a homography. Think of the images as nodes in a graph and the image pairs as edges. The images that will form the mosaic are the connected components. (If for some reason there is more than one connected component, pick the largest.) Second, you will need to pick an “anchor” image that will remain fixed while the other images are mapped onto it. This should in some sense be the “center” of the set of images in the connected component. Third, you need to compute the transformations that map the images onto this anchor image. This can get tricky quite quickly, so please do something very easy using only the results of matching pairs of images. In particular, if I_0 is the anchor and I_i is successfully matched with I_0 , then use the transformation homography computed between them. If I_i does not have a homography with I_0 , but there is another image I_j that does, then “compose” the transformations: I_i onto I_j onto I_0 . This is not as hard as it might sound. In particular, if $\mathbf{H}_{j,i}$ is the estimated transformation matrix from I_i onto I_j and if $\mathbf{H}_{0,j}$ is the estimated transformation matrix from I_j onto I_0 , then $\mathbf{H}_{0,i} = \mathbf{H}_{0,j}\mathbf{H}_{j,i}$ is a good estimate of the transformation from I_i onto I_0 . In the data provided there **will not** be any cases where you need to compose more than two transformations if you choose the anchor correctly. Note that commercial software that builds multi-image mosaics uses much more sophisticated methods to estimate $\mathbf{H}_{0,i}$.

Command Line and Output

Your program should run with the following very simple command line:

```
python hw3_align.py in_dir out_dir
```

where `in_dir` is the path to the directory containing input images. We will run some of your submissions to test them. The code should write all images to `out_dir`, which should be a different directory from `in_dir`. This will avoid clutter across multiple runs. Your code will need to output (via `print` statements) a significant amount of text as described above. For each mosaic you create, make the file name be the composition of the names of the input file prefixes, in sorted order. For example, if the images are `bar.jpg`, `cheese.jpg` and `foo.jpg`, then the mosaic of the first two will be `bar_cheese.jpg` and the mosaic of all three will be `bar_cheese_foo.jpg`. Use the extension from the first image (all images will be jpg or JPG). Note that for image pairs that do form mosaics, there will be four output images — the images that result from steps 2, 4, 6 and 8. For pairs that do not form mosaics, there will be fewer output images, depending on which decision (steps 3, 5 or 7) stopped the computation. There will always be an output image from step 2.

Write Up and Code

Please generate a write-up describing your algorithms, your decision criteria, your blending algorithms, and your overall results. Evaluate **both strengths and weaknesses**, using image pairs — perhaps including some we did not provide — to illustrate the stages of your computation and your insights.

One requirement is to make a table summarizing the results on all the image pairs, including the matching results, the number and percentage of inliers to F and to H (if they were estimated), and the final decision your algorithm made. This will take some time to generate but it is the type of analysis you will need to learn how to make to evaluate computer vision and machine learning algorithms. You don't have to make this beautiful, just make it clear and easy to follow.

The actual text should be no more than a page or so, single-spaced, but the document should be longer because of the results table and the illustrating images.

Finally, make sure your code is clean, reasonably efficient, documented, and well-structured.

Complete Submission

Your final submission will be a single zip file that includes the following:

1. Your .py file — no notebook for this!
2. The text output files from running your code on **each** of the image sets provided, plus other image sets you'd like to show. One additional suggest is to run your algorithm on two images taken from different sets.
3. As many image results as you need to illustrate your successes (and failures), both in forming mosaics and in deciding not to do so!
4. Your final write-up.

The zip file will be limited to 60MB. This means it is unlikely that you can include all image results.